

Project - 2

```
[12]: import numpy as np
import pandas as pd
import re
import matplotlib.pyplot as plt
from matplotlib import style
%matplotlib inline
```

```
[13]: #Importing all 3 Datasets
users_data = pd.read_csv("users.dat", sep="::", header=None,
    ↳ names=['UserID', 'Gender', 'Age', 'Occupation', 'Zip-code'],
        dtype={'UserID': np.int32, 'Gender': np.str, 'Age': np.int32,
    ↳ 'Occupation': np.int32, 'Zip-code': np.str},
        engine='python')
movie_data = pd.read_csv("movies.dat",
    ↳ sep="::", header=None,
    ↳ names=['MovieID', 'Title', 'Genres'],
        dtype={'MovieID': np.int32, 'Title': np.str, 'Genres':
    ↳ np.str}, engine='python')
ratings_data = pd.read_csv("ratings.dat",
    ↳ sep="::", header=None,
    ↳ names=['UserID', 'MovieID', 'Rating', 'Timestamp'],
        dtype={'UserID': np.int32, 'MovieID': np.int32, 'Rating': np.
    ↳ int32, 'Timestamp': np.str}, engine='python')
```

```
[3]: #Analysing the Datasets
#1) Users Data
users_data.head()
```

```
[3]:   UserID Gender  Age  Occupation Zip-code
0      1      F   1      10      48067
1      2      M  56      16      70072
2      3      M  25      15      55117
3      4      M  45       7      02460
4      5      M  25      20      55455
```

```
[4]: users_data.isnull().sum()
```

```
[4]: UserID      0
Gender      0
Age      0
Occupation 0 Zip-
code 0
dtype: int64
```

```
[5]: users_data.shape
```

```
[5]: (6040, 5)
```

```
[8]: #2) Movie Data
movie_data.head()
```

```
[8]:      MovieID      Title      Genres
0         1  Toy Story (1995) Animation|Children's|Comedy
1         2  Jumanji (1995)  Adventure|Children's|Fantasy
2         3  Grumpier Old Men (1995)  Comedy|Romance
3         4  Waiting to Exhale (1995)  Comedy|Drama
4         5  Father of the Bride Part II (1995)  Comedy
```

```
[9]: movie_data.isnull().sum()
```

```
[9]: MovieID      0
     Title      0
     Genres      0
     dtype: int64
```

```
[10]: movie_data.shape
```

```
[10]: (3883, 3)
```

```
[11]: #3) Rating data
ratings_data.head()
```

```
[11]:      UserID  MovieID  Rating  Timestamp
0         1         1193        5  978300760
1         1          661        3  978302109
2         1          914        3  978301968
3         1         3408        4  978300275
4         1         2355        5  978824291
```

```
[12]: ratings_data.isnull().sum()
```

```
[12]: UserID 0  MovieID
0
Rating      0
Timestamp    0
     dtype: int64
```

```
[13]: ratings_data.shape
```

```
[13]: (1000209, 4)
```

```
[14]: #Merging the Dataset and creating a Master Dataset
#Merging Users dataset and ratings dataset
Master_Data = pd.merge(users_data,ratings_data,on = 'UserID')
Master_Data.head()
```

```
[14]:      UserID  Gender  Age  Occupation  Zip-code  MovieID  Rating  Timestamp
```

0	1	F	1	10	48067	1193	5	978300760
1	1	F	1	10	48067	661	3	978302109
2	1	F	1	10	48067	914	3	978301968
3	1	F	1	10	48067	3408	4	978300275
4	1	F	1	10	48067	2355	5	978824291

```
[15]: #Merging Master Dataset and movie dataset
Master_Data=pd.merge(Master_Data,movie_data,on = 'MovieID')
Master_Data.head()
```

```
[15]:  UserID  Gender  Age  Occupation  Zip-code  MovieID  Rating  Timestamp \
0      1      F      1      10      48067  1193      5      978300760
1      2      M     56     16     70072  1193      5      978298413
2     12      M     25     12     32793  1193      4      978220179
3     15      M     25      7     22903  1193      4      978199279
4     17      M     50      1     95350  1193      5      978158471
```

	Title	Genres
0	One Flew Over the Cuckoo's Nest (1975)	Drama
1	One Flew Over the Cuckoo's Nest (1975)	Drama
2	One Flew Over the Cuckoo's Nest (1975)	Drama
3	One Flew Over the Cuckoo's Nest (1975)	Drama
4	One Flew Over the Cuckoo's Nest (1975)	Drama

```
[16]: #Preparing the Master dataset as required
Master_Data = Master_Data.drop(['Zip-code'],axis=1)
Master_Data = Master_Data.drop(['Timestamp'],axis=1)
```

```
[17]: Master_Data = _
      _Master_Data[['UserID', 'Gender', 'Age', 'Occupation', 'MovieID', 'Title', 'Genres',
      'Rating']] Master_Data.head()
```

```
[17]:  UserID  Gender  Age  Occupation  MovieID \
0      1      F      1      10      1193
1      2      M     56     16      1193
2     12      M     25     12     1193  3      15      M     25      7      1193
4     17      M     50      1      1193
```

	Title	Genres
Rating 0	One Flew Over the Cuckoo's Nest (1975)	Drama 5
1	One Flew Over the Cuckoo's Nest (1975)	Drama 5
2	One Flew Over the Cuckoo's Nest (1975)	Drama 4

```
3 One Flew Over the Cuckoo's Nest (1975) Drama 4
4 One Flew Over the Cuckoo's Nest (1975) Drama 5
```

```
[57]: #Data Visualizations
      #1) User Age Distribution
```

```
[18]: Age_count = users_data['Age'].value_counts()
      Age_count
```

```
[18]: 25    2096
      35    1193
      18    1103
      45     550
      50     496
      56     380
      1      222
      Name: Age, dtype: int64
```

```
[19]: Age_Category = ('Under 18', '18-24', '25-34', '35-44', '45-49', '50-55', '56+')
      x_position = np.arange(len(Age_Category))
      x_position
```

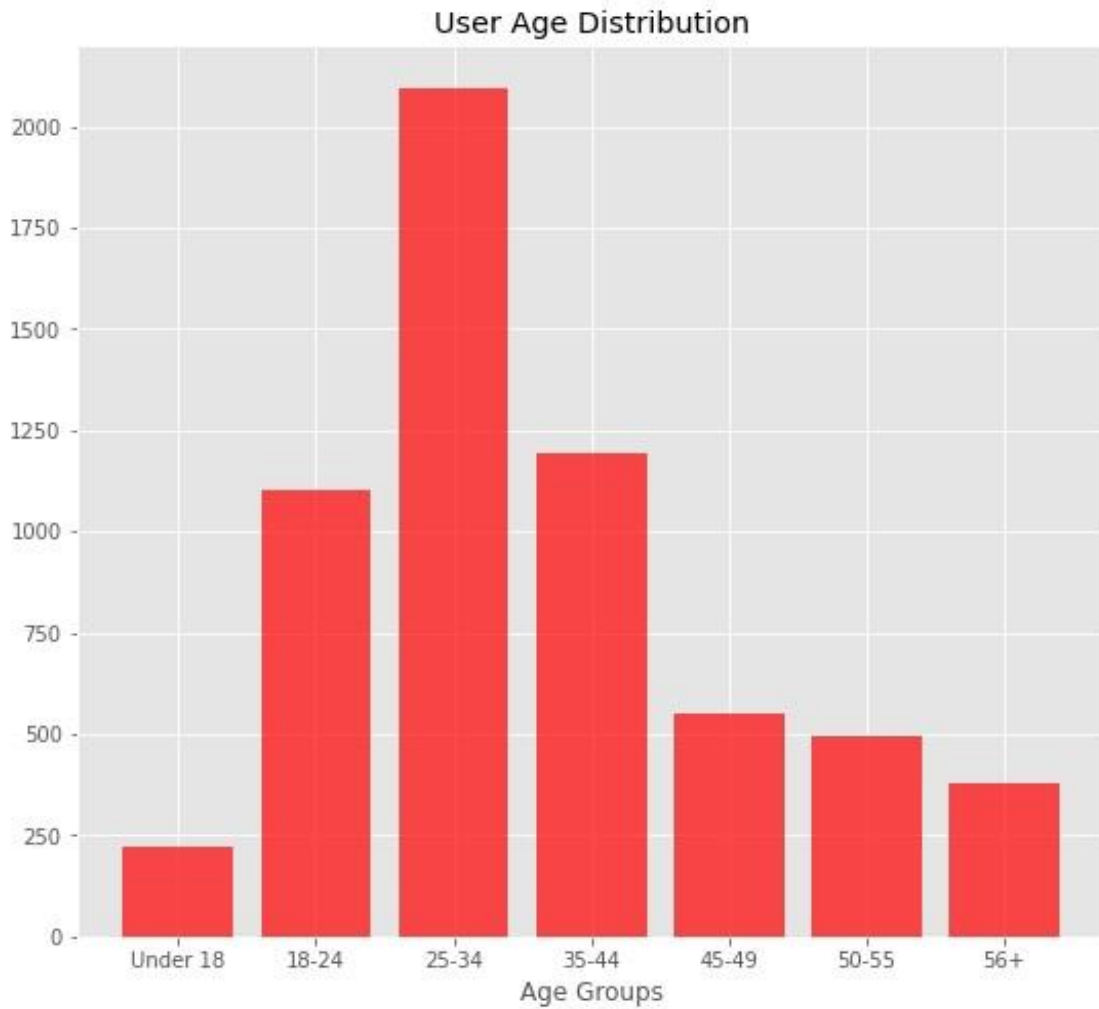
```
[19]: array([0, 1, 2, 3, 4, 5, 6])
```

```
[20]: Age_Values = _
      → [Age_count[1],Age_count[18],Age_count[25],Age_count[35],Age_count[45],Age_count[50],Age_count[56]]
      Age_Values
```

```
[20]: [222, 1103, 2096, 1193, 550, 496, 380]
```

```
[10]: #plotting bar chart
      style.use('ggplot')
      plt.figure(figsize=(9,8))
      plt.bar(x_position, Age_Values, align='center', color='r', alpha=0.7)
      #set the y axis lable
      plt.xlabel('Age Groups')
      #set the bar value
      plt.xticks(x_position, Age_Category)
```

```
#set the title
plt.title('User Age Distribution')
plt.show()
```



```
[142]: #The above age distribution shows that most of the users are 25-34 years old
```

```
[1]: #2) User ratings of movie Toy Story
```

```
[21]: #Fetching thr Movie ID of Toy Story
movie_data.MovieID[movie_data.Title=='Toy Story (1995)']
```

```
[21]: 0    1
      Name: MovieID, dtype: int32
```

```
[22]: toystory_data = ratings_data[ratings_data.MovieID==1]
toystory_data.head(10)
```

```
[22]:   UserID MovieID Rating Timestamp
40      1      1      5  978824268
469     6      1      4  978237008
581     8      1      4  978233496
711     9      1      5  978225952
837    10      1      5  978226474
1966   18      1      4  978154768
2276   19      1      5  978555994
2530   21      1      3  978139347
2870   23      1      4  978463614
3405   26      1      3  978130703
```

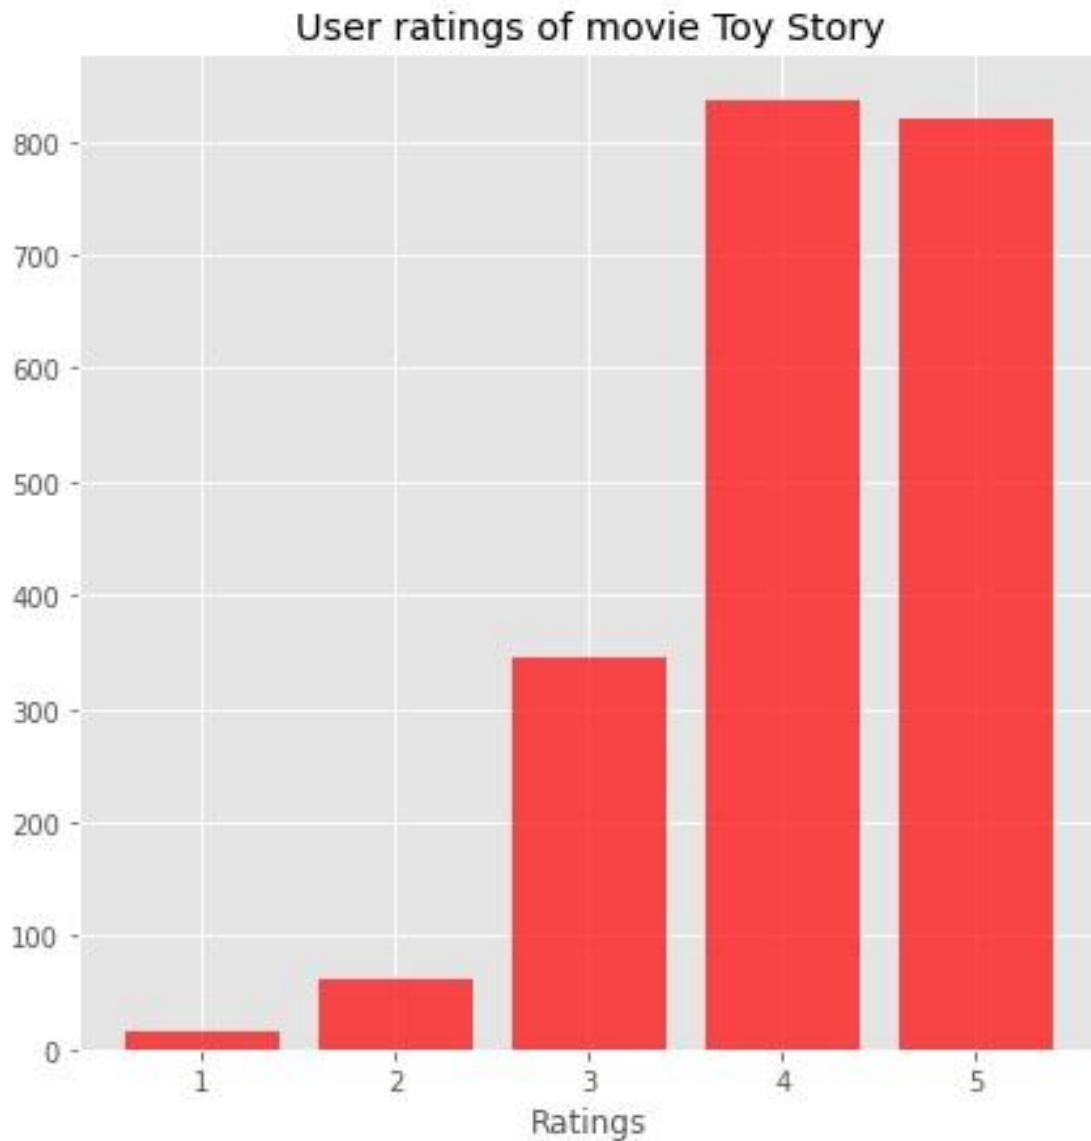
```
[23]: movie_ratings_toystory = toystory_data.groupby('Rating').size()
movie_ratings_toystory
```

```
[23]: Rating
1      16
2      61
3     345
4     835
5     820
dtype: int64
```

```
[24]: ratings_type = ('1','2','3','4','5')
x_pos = np.arange(len(ratings_type))
x_pos
```

```
[24]: array([0, 1, 2, 3, 4])
```

```
[15]: #plotting bar chart
style.use('ggplot')
plt.figure(figsize=(7,7))
plt.bar(x_pos,movie_ratings_toystory,align='center',color='r',alpha=0.7)
#set the y axis lable
plt.xlabel('Ratings')
#set the bar value
plt.xticks(x_pos,ratings_type)
#set the title
plt.title('User ratings of movie Toy Story ')
plt.show()
```



```
[53]: #The above plot shows that the movie 'Toystory' has got 4 **  
(stars) maximum
```

```
[20]: #3) Top 25 movies by viewership rating
```

```
[25]: #Fetching the Data and ratings of each movie by aggrega  
movie_rating = Master_Data.groupby(['Title'],  
as_index=False) average_movie_ratings =  
movie_rating.agg({'Rating': 'mean'})  
average_movie_ratings.head(25)
```

```
[25]:
```

	Title	Rating
0	\$1,000,000 Duck (1971)	3.027027
1	'Night Mother (1986)	3.371429
2	'Til There Was You (1997)	2.692308
	3 'burbs, The (1989)	2.910891
4	...And Justice for All (1979)	3.713568
5	1-900 (1994)	2.500000
6	10 Things I Hate About You (1999)	3.422857
7	101 Dalmatians (1961)	3.596460
8	101 Dalmatians (1996)	3.046703
9	12 Angry Men (1957)	4.295455
10	13th Warrior, The (1999)	3.158667
11	187 (1997)	2.745455
12	2 Days in the Valley (1996)	3.283217
13	20 Dates (1998)	2.856115
14	20,000 Leagues Under the Sea (1954)	3.702609
15	200 Cigarettes (1999)	2.883978
16	2001: A Space Odyssey (1968)	4.068765
17	2010 (1984)	3.417021
18	24 7: Twenty Four Seven (1997)	4.000000
19	24-hour Woman (1998)	1.777778
20	28 Days (2000)	3.065347
21	3 Ninjas: High Noon On Mega Mountain (1998)	1.361702
22	3 Strikes (2000)	2.750000
23	301, 302 (1995)	2.888889
24	39 Steps, The (1935)	4.075099

```
[26]: top_25_movies = average_movie_ratings.sort_values('Rating', ascending=False).
      → head(25)
      top_25_movies
```

```
[26]:
```


	Title	Rating
3477	Ulysses (Ulissee) (1954)	5.000000
2025	Lured (1947)	5.000000
1203	Follow the Bitch (1998)	5.000000
407	Bittersweet Motel (2000)	5.000000
3087	Song of Freedom (1936)	5.000000
2453	One Little Indian (1973)	5.000000
3044	Smashing Time (1967)	5.000000
2903	Schlafes Bruder (Brother of Sleep) (1995)	5.000000
1297	Gate of Heavenly Peace, The (1995)	5.000000
249	Baby, The (1973)	5.000000
1622	I Am Cuba (Soy Cuba/Ya Kuba) (1964)	4.800000
1870	Lamerica (1994)	4.750000
199	Apple, The (Sib) (1998)	4.666667
2883	Sanjuro (1962)	4.608696

2940Seven Samurai (The Magnificent Seven) (Shichin...4.560510
 2970 Shawshank Redemption, The (1994)4.554558
 1354 Godfather, The (1972)4.524966
 713 Close Shave, A (1995)4.520548
 3504 Usual Suspects, The (1995) 4.517106 2901 Schindler's
 List (1993) 4.510417
 3675 Wrong Trousers, The (1993)4.507937
 996 Dry Cleaning (Nettoyage sec) (1997)4.500000
 1690 Inheritors, The (Die Siebtelbauern) (1998)4.500000
 2056 Mamma Roma (1962)4.500000

342

Bells, The (1926) 4.500000

```
[91]: top_25_movies.plot(kind='barh',alpha=0.6,figsize=(7,7))
plt.xlabel("Viewership Ratings Count")
plt.ylabel("Movies (Top 25) ")
plt.title("Top 25 movies by viewership rating ")
plt.show()
```




```
[27]: #Ratings for all the movies reviewed by for a particular user of
user id = 2696 user_rating_data =
Master_Data[Master_Data['UserID']==2496] user_rating_data =
user_rating_data[['UserID', 'MovieID', 'Title', 'Rating']]
user_rating_data.head(10)
```

```
[27]:   UserID MovieID Title Rating
668    2496   1193 One Flew Over the Cuckoo's Nest (1975) 5
2518    2496    914 My Fair Lady (1964) 4
```

8506	2496	1287	Ben-Hur (1959)	5
9492	2496	2804	Christmas Story, A (1983)	4
14173	2496	2398	Miracle on 34th Street (1947)	5
16319	2496	1035	Sound of Music, The (1965)	5
17581	2496	2791	Airplane! (1980)	3
19798	2496	3105	Awakenings (1990)	3
24263	2496	1270	Back to the Future (1985)	2
26818	2496	527	Schindler's List (1993)	5

```
[96]: # plotting the above data
plt.scatter(x=user_rating_data['MovieID'].head(20), _
→ y=user_rating_data['Rating'].head(20))
plt.show()
```



```
[19]: #Feature Engineering
# 1) Find out all the unique genres
```

```
[28]: genres = Master_Data['Genres'].str.split("|")
genres
```

```
[28]: 0      [Drama]
      1      [Drama]
      2      [Drama]
      3      [Drama]
```

```

4                                [Drama]
...
1000204                        [Documentary]
1000205                        [Drama]
1000206                        [Drama]
1000207                        [Comedy, Drama, Western]
1000208                        [Documentary]
Name: Genres, Length: 1000209, dtype: object

```

```

[29]: unique_genres = set()
      for gen in genres:
          unique_genres = unique_genres.union(set(gen))

```

```

[33]: unique_genres

```

```

[33]: {'Action',
      'Adventure',
      'Animation',
      "Children's",
      'Comedy',
      'Crime',
      'Documentary',
      'Drama',
      'Fantasy',
      'Film-Noir',
      'Horror',
      'Musical',
      'Mystery',
      'Romance',
      'Sci-Fi',
      'Thriller',
      'War',
      'Western'}

```

```

[35]: # 2) Create a separate column for each genre category with a one-hot encoding (
      → 1 and 0)

```

```

[30]: oneHotGenre = Master_Data["Genres"].str.get_dummies("|")
      oneHotGenre.head()

```

```

[30]: Action Adventure Animation Children's Comedy Crime Documentary \
0      0  0  0  0  0  0  0  0
1      0  0  0  0  0  0  0  0
2      0  0  0  0  0  0  0  0
3      0  0  0  0  0  0  0  0
4      0  0  0  0  0  0  0  0

      Drama Fantasy Film-Noir Horror Musical Mystery Romance Sci-Fi \

```

0	1	0	0	0	0	0	0	0
1	1	0	0	0	0	0	0	0
2	1	0	0	0	0	0	0	0
3	1	0	0	0	0	0	0	0
4	1	0	0	0	0	0	0	0

	Thriller	War	Western
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

```
[31]: oneHotGenre = pd.concat([Master_Data,oneHotGenre],axis=1)
oneHotGenre.head()
```

```
[31]:  UserID  Gender  Age  Occupation  MovieID  \
0         1    F     1      10      1193
1         2    M    56      16      1193
2        12    M    25      12      1193  3      15    M    25    7      1193
4        17    M    50           1      1193
```

	Title	Genres	Rating	Action	Adventure	\
0	One Flew Over the Cuckoo's Nest (1975)	Drama	5	0	0	
1	One Flew Over the Cuckoo's Nest (1975)	Drama	5	0	0	
2	One Flew Over the Cuckoo's Nest (1975)	Drama	4	0	0	
3	One Flew Over the Cuckoo's Nest (1975)	Drama	4	0	0	
4	One Flew Over the Cuckoo's Nest (1975)	Drama	5	0	0	

	...	Fantasy	Film-Noir	Horror	Musical	Mystery	Romance	Sci-Fi	\
0	...	0	0	0	0	0	0	0	
1	...	0	0	0	0	0	0	0	
2	...	0	0	0	0	0	0	0	
3	...	0	0	0	0	0	0	0	
4	...	0	0	0	0	0	0	0	

	Thriller	War	Western
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

[5 rows x 26 columns]

```
[49]: oneHotGenre.columns
```

```
[49]: Index(['UserID', 'Gender', 'Age', 'Occupation', 'MovieID',
'Title', 'Genres', 'Rating', 'UserID', 'Gender', 'Age', 'Occupation',
'MovieID', 'Title',
'Genres', 'Rating', 'UserID', 'Gender', 'Age', 'Occupation',
'MovieID',
'Title', 'Genres', 'Rating', 'UserID', 'Gender', 'Age',
'Occupation',
'MovieID', 'Title', 'Genres', 'Rating', 'UserID', 'Gender',
'Age',
'Occupation', 'MovieID', 'Title', 'Genres', 'Rating',
'Action',
'Adventure', 'Animation', 'Children's', 'Comedy', 'Crime',
'Documentary', 'Drama', 'Fantasy', 'Film-Noir', 'Horror',
'Musical',
'Mystery', 'Romance', 'Sci-Fi', 'Thriller', 'War', 'Western'],
dtype='object')
```

```
[17]: # 3) Determine the features affecting the ratings of any particular movie
```

```
[32]: Features_Data =Master_Data.copy()
Features_Data
```

```
[32]:
```

	UserID	Gender	Age	Occupation	MovieID \
0	1	F	1	10	1193
1	2	M	56	16	1193
2	12	M	25	12	1193
3	15	M	25	7	1193
4	17	M	50	1	1193
...
1000204	5949	M	18	17	2198
1000205	5675	M	35	14	2703
1000206	5780	M	18	17	2845
1000207	5851	F	18	20	3607
1000208	5938	M	25	1	2909

	Title	Genres \
0	One Flew Over the Cuckoo's Nest (1975)	Drama
1	One Flew Over the Cuckoo's Nest (1975)	Drama
2	One Flew Over the Cuckoo's Nest (1975)	Drama
3	One Flew Over the Cuckoo's Nest (1975)	Drama
4	One Flew Over the Cuckoo's Nest (1975)	Drama
...
1000204	Modulations (1998)	Documentary
1000205	Broken Vessels (1998)	Drama
1000206	White Boys (1999)	Drama
1000207	One Little Indian (1973)	Comedy Drama Western

```
1000208                    Five Wives, Three Secretaries and Me (1998)
                             Documentary
```

```

      Rating
0          5
1          5
2          4
3          4
4          5
...
1000204    5
1000205    3
1000206    1
1000207    5
1000208    4
```

```
[1000209 rows x 8 columns]
```

```
[35]: #Fetching the year which the movie was released
Features_Data[["Title", "Year"]] =
    Features_Data.Title.str.extract("(.)\s\((. \d+)\)", expand=True)
Features_Data = Features_Data.drop(['Title'], axis=1)
Features_Data
```

```
[35]:
      UserID Gender Age Occupation MovieID      Genres \
0         1     F    1      10      1193  Drama 1     2     M     56
      16      1193  Drama
2         12     M    25      12      1193  Drama
3         15     M    25       7      1193  Drama
4         17     M    50       1      1193  Drama
...
1000204    5949  M    18      17      2198  Documentary 1000205
      5675  M    35      14      2703  Drama
1000206    5780  M    18      17      2845  Drama
1000207    5851  F    18      20      3607  Comedy|Drama|Western
1000208    5938  M    25       1      2909  Documentary
```

```

      Rating Year
0          5  1975
1          5  1975
2          4  1975
3          4  1975
4          5  1975
...
1000204    5  1998
```

```

1000205      3  1998
1000206      1  1999
1000207      5  1973
1000208      4  1998
[1000209 rows x 8 columns]

```

```

[36]: #Calculating the age of movies
Features_Data['Year'] = Features_Data.Year.astype(int)
Features_Data['Movie_Age'] = 2000 -Features_Data['Year']
Features_Data

```

```

[36]:      UserID Gender Age Occupation MovieID      Genres \
0         1      F    1         10      1193 Drama 1      2      M      56
      16      1193 Drama
2         12      M    25         12      1193 Drama
3         15      M    25          7      1193 Drama
4         17      M    50          1      1193 Drama
...
1000204      5949 M    18         17      2198 Documentary 1000205
      5675 M    35         14      2703 Drama
1000206      5780 M    18         17      2845 Drama
1000207      5851 F    18         20      3607 Comedy|Drama|Western
1000208      5938 M    25          1      2909 Documentary

      Rating Year Movie_Age
0         5  1975  25
1         5  1975  25
2         4  1975  25
3         4  1975  25
4         5  1975  25
...
1000204      5  1998  2
1000205      3  1998  2
1000206      1  1999  1 1000207  5      1973  27
1000208      4  1998      2

[1000209 rows x 9 columns]

```

```

[37]: #Creating Gender variable as integer type
Features_Data['Gender'] = Features_Data.Gender.replace('F',1)
Features_Data['Gender'] = Features_Data.Gender.replace('M',0)
Features_Data['Gender'] = Features_Data.Gender.astype(int)
Features_Data.head()

```

```

[37]:      UserID Gender Age Occupation MovieID      Rating Year Movie_Age
      Genres
0         1      1    1         10      1193 Drama      5  1975      25

```



1	2	0	56	16	1193	Drama	5	1975	25
2	12	0	25	12	1193	Drama	4	1975	25
3	15	0	25	7	1193	Drama	4	1975	25
4	17	0	50	1	1193	Drama	5	1975	25

```
[38]: #Checking the correlation of features with Rating
Features_Data[['Gender','Occupation', 'Age', 'Movie_Age']].
->corrwith(Features_Data['Rating'])
```

```
[38]: Gender      0.019861
Occupation  0.006753
Age         0.056869
Movie_Age   0.156946
dtype: float64
```


```
[103]: #Movie_Age has the most positive relationship with Rating
```

```
[105]: #Occupation relationship with Rating
Features_Data.groupby(["Occupation","Rating"]).size().unstack().
->plot(kind='bar',stacked=False,legend=True)
plt.show()
```




```
[107]: #Gender relationship with Rating
#1 -> Male, 0 -> Female
```

```
Features_Data.groupby(["Gender", "Rating"]).size().unstack().
    .plot(kind='bar', stacked=False, legend=True)
plt.show()
```




[108]: *#Age relationship with Rating*

```
Features_Data.groupby(["Age", "Rating"]).size().unstack(
    ).
    .plot(kind='bar', stacked=False, legend=True)
plt.show()
```



```
[109]: #Movie_Age relationship with Rating
Features_Data.groupby(["Movie_Age","Rating"]).size().unstack().
→plot(kind='bar',stacked=False,legend=True)
plt.show()
```



```
[39]: #To Predict the values of rating we are using Logistic regression
```

```
[75]: # Assign independent variables to X dataset
X = Master_Data[['Age','Occupation','MovieID']].head(500)
X
```

```
[75]:
```

	Age	Occupation	MovieID
0	1	10	1193
1	56	16	1193
2	25	12	1193
3	25	7	1193
4	50	1	1193
..
495	25	2	1193
496	18	4	1193
497	25	12	1193

```
498  18          4    1193
```

```
499  45          14    1193
```

```
[500 rows x 3 columns]
```

```
[76]: # Assign dependent variables to Y dataset
Y = Master_Data['Rating'].head(500)
Y
```

```
[76]: 0      5
```

```
1      5
```

```
2      4
```

```
3      4
```

```
4      5
```

```
..
```

```
495    4
```

```
496    5
```

```
497    5
```

```
498    5
```

```
499    5
```

```
Name: Rating, Length: 500, dtype: int32
```

```
[77]: # view the shape for both axes
print (X.shape)
print (Y.shape)
```

```
(500, 3)
```

```
(500,)
```

```
[78]: # Splitting the data into training & testing
datasets(70:30) import sklearn from
sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = sklearn.model_selection.
    ,train_test_split(X,Y,random_state=2,test_size=0.3)
```

```
[79]: # use the Logistic regression estimator
from sklearn.linear_model import
LogisticRegression logReg =
LogisticRegression()
```

```
[81]: # fit data into the Logistic regression estimator
logReg.fit(X_train,Y_train)
```

```
/usr/local/lib/python3.7/site-  
packages/sklearn/linear_model/_logistic.py:940:  
ConvergenceWarning: lbfgs failed to converge  
(status=1): STOP: TOTAL NO. of ITERATIONS REACHED  
LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as
shown in: [https://scikit-
learn.org/stable/modules/preprocessing.html](https://scikit-learn.org/stable/modules/preprocessing.html)

Please also refer to the documentation for alternative solver
options:

```
https://scikit-  
learn.org/stable/modules/linear\_model.html#logisticregression  
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

```
[81]: LogisticRegression(C=1.0, class_weight=None, dual=False,  
        fit_intercept=True, intercept_scaling=1,  
        l1_ratio=None, max_iter=100, multi_class='auto',  
        n_jobs=None, penalty='l2', random_state=None,  
        solver='lbfgs', tol=0.0001, verbose=0,  
        warm_start=False)
```

```
[82]: #Model Evaluation  
# predict the outcome using Logistic regression estimator  
y_predict=logReg.predict(X_test)
```

```
[83]: y_predict
```

```
[83]: array([5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5,  
5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5,  
5, 5, 5,  
5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5,  
5,  
5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5,  
5,  
5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5,  
5,  
5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5,  
5,  
5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5,  
5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5],  
dtype=int32)
```

```
[84]: # Calculate the accuracy of the  
model from sklearn.metrics import  
accuracy_score  
accuracy_score(y_predict,Y_test)
```

```
[84]: 0.5866666666666667
```

```
[85]: #Check model performance on new dataset
# create Example object with new values for prediction
X_new = [[25,7,1193],[18,17,2198]]
```

```
[86]: logReg.predict(X_new)
```

```
[86]: array([5, 5], dtype=int32)
```

```
[89]: from sklearn import metrics print
(metrics.confusion_matrix(Y_test, y_predict))
print (metrics.classification_report(Y_test,
y_predict))
```

```
[[ 0  0  0  0  1]
 [ 0  0  0  0  2]
 [ 0  0  0  0  9]
 [ 0  0  0  0 50]
 [ 0  0  0  0 88]]
```

```
precision recall f1-score support
```

```
10.00 0.00 0.00 1 2 0.00 0.00
0.00 2 3 0.00 0.00 0.00 9
4 0.00 0.00 0.00 50
5 0.59 1.00 0.74 88
```

```
accuracy 0.59 150 macro avg 0.12 0.20
0.15 150 weighted avg 0.34 0.59 0.43 150
```

```
/usr/local/lib/python3.7/site-
packages/sklearn/metrics/_classification.py:1272:
UndefinedMetricWarning: Precision and F-score are ill-defined and
being set to 0.0 in labels with no predicted samples. Use
`zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
```

```
[ ]:
```